

# HTTP/2 and gRPC

© Riccardo Sisto, Politecnico di Torino, 2020-2024

References for study:


- Ilya Grigorik, "High Performance browser Networking", chapter 12 (<https://hpbnn.co/http2/>)
- gRPC Tutorials (<https://grpc.io/docs/tutorials/>)
- Protocol Buffers tutorials (<https://protobuf.dev/programming-guides/>)

# HTTP/2 and HTTP/3


- Evolutions of HTTP for better efficiency/performance
  - HTTP/2: optimized for networks with low error rates (e.g., within data centers)
  - HTTP/3: optimized to reduce latency for long distance or high error rate connections
- Same semantics as HTTP/1.1, but different data encoding and use of transport protocols
  - No impact on existing REST APIs => they can be ported to HTTP/2 and HTTP/3 smoothly

# TCP/HTTP1.1 Performance Issues

- High latency for establishing connections



**HTTP: 2RTT**



**HTTPS: 3RTT**

- Multiplexing requires opening multiple TCP connections
- Character encoding is not very efficient

# HTTP/2

- **HTTP/2:** Initially driven by Google (SPDY), later on taken over by IETF (RFC 7540)

	HTTP/2	HTTP/1.1
data encoding	binary	character-oriented
multiplexing	on single TCP connection	on multiple TCP connections
compression	header and body	only body
push possibility	yes	no


# HTTP/3

- RFC 9114 (June 2022)
- TCP replaced by QUIC+UDP


HTTP/1 HTTP/2
(TLS)
TCP
IP

HTTP/3
QUIC (+TLS)
UDP
IP

**HTTP/1 HTTP/2**



**HTTP/3**



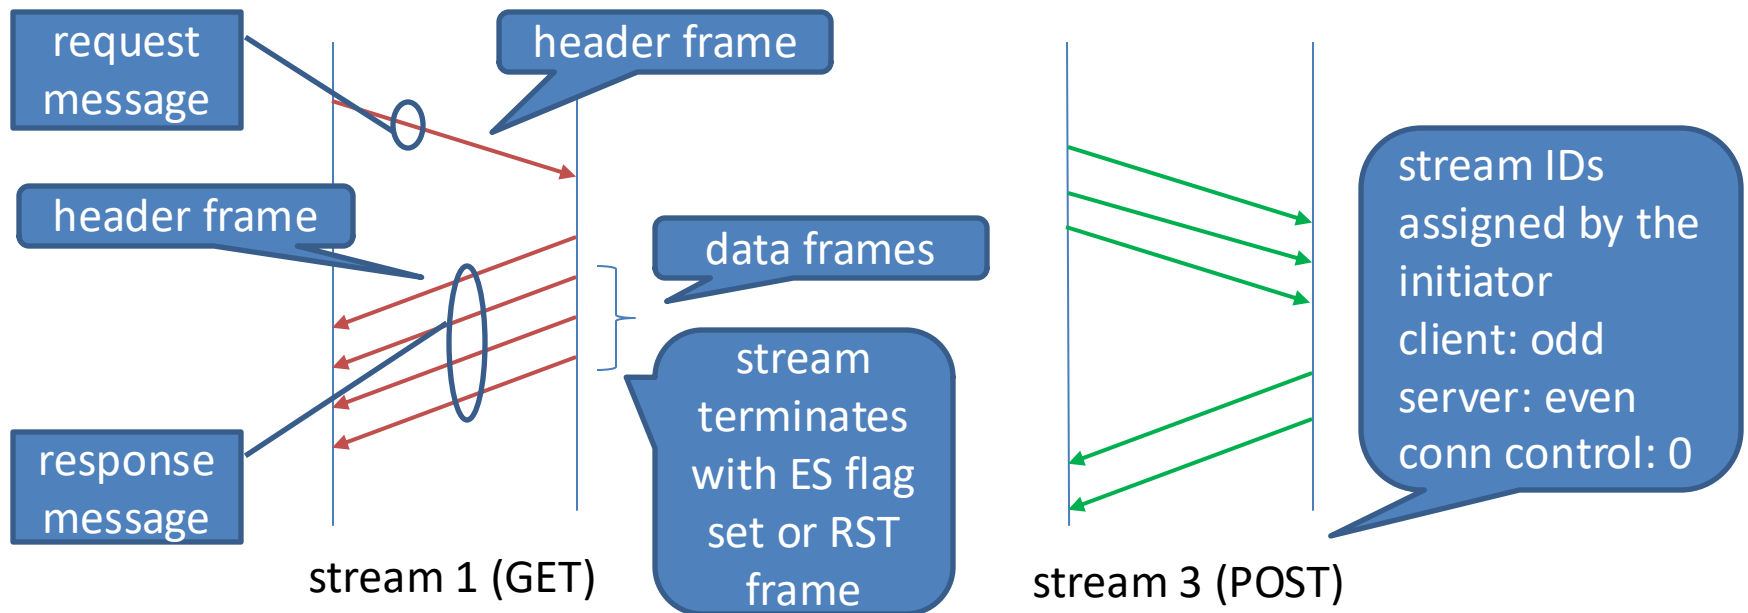
# HTTP/2 HTTP/3

## Adoption and Performance


- Adoption in 2021 (Web Almanac, Percentage of Page loads):
  - HTTP/1 40%
  - HTTP/2 45% (but push adoption <1%)
  - HTTP/3 15%
- Performance:
  - HTTP/2: big boost wrt HTTP/1
  - HTTP/3: better than HTTP/2 only for long distance or high error rate/unreliable connections

# HTTP/2: Frames and Streams

- Frame: binary message exchanged by client and server
- Stream: independent, bidirectional sequence of frames
  - corresponds to an ordered sequence of logical operations (Request/Response, Push), i.e., of HTTP messages




# Frame Layout





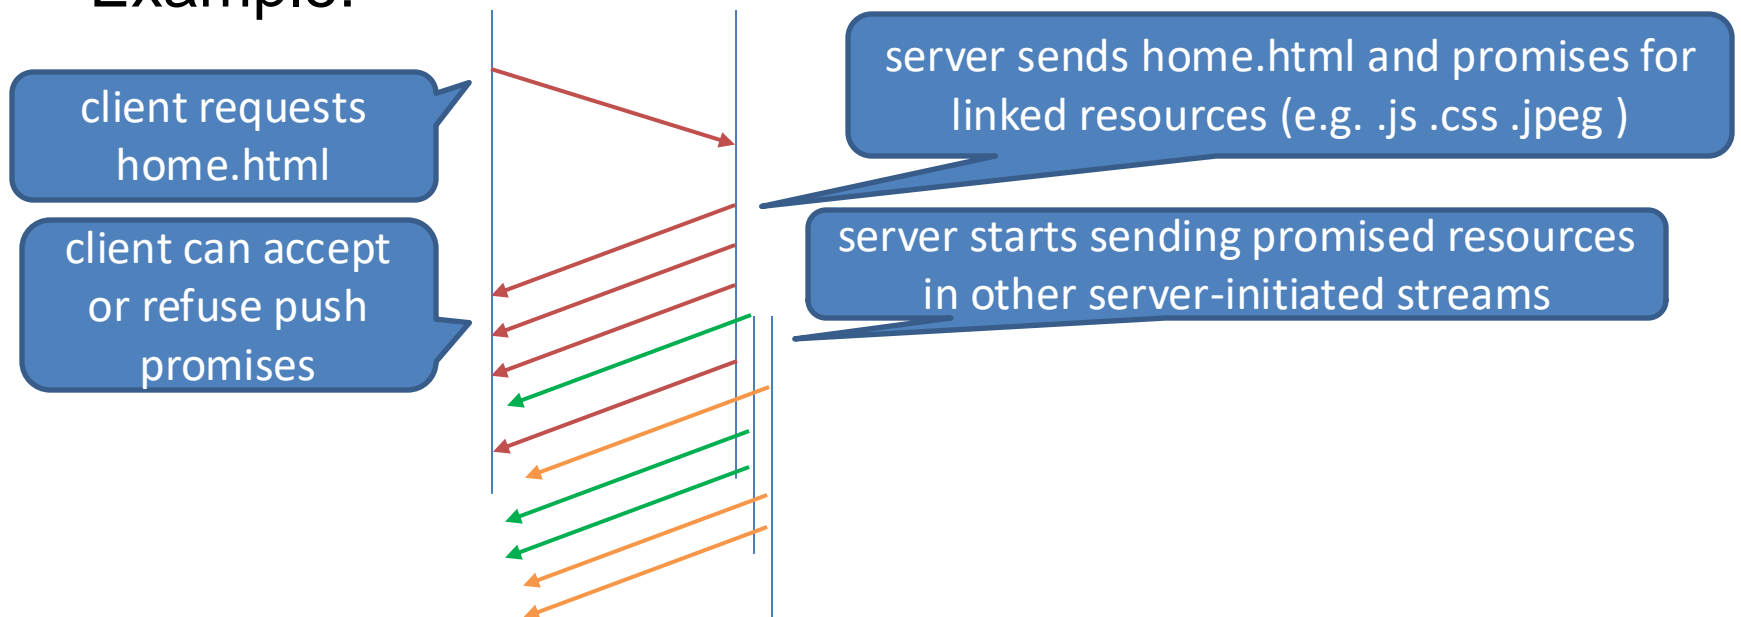
# Stream Multiplexing

- Streams are multiplexed on a single TCP connection
  - ordering of the frames of each stream preserved
  - window-based flow control added to regulate stream rates





# Push Mechanism

- A server can push additional representations, anticipating client's requests
  - a way to trade bandwidth for latency
- Example:

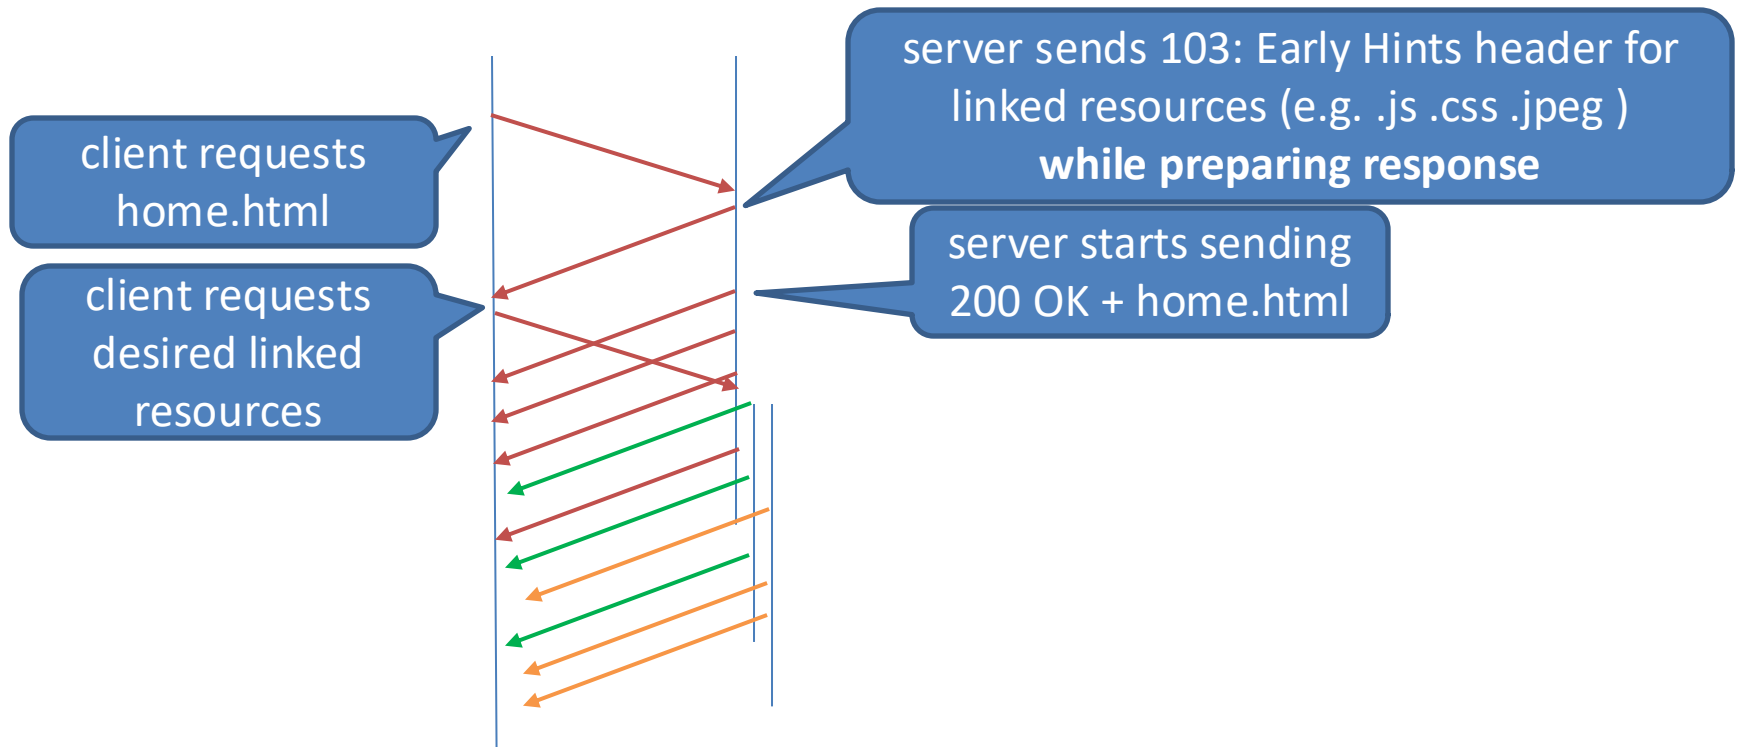


# Push vs No-Push



# Early Hints

- Alternative to server push which saves bandwidth
- Example:



# Main Frame Types

- HEADERS
- DATA
- SETTINGS
- RST\_STREAM                      signal termination of a stream
- PUSH\_PROMISE
- GOAWAY                      initiate (graceful) connection termination
- WINDOW\_UPDATE              flow control message

# Initiating an HTTP/2 Connection

- If the client knows the server supports HTTP/2
  - After initiating the TCP connection, the client sends the 24-bytes preface string ("PRI \* HTTP/2.0\r\n\r\nSM\r\n\r\n") followed by a SETTINGS frame
  - The server responds with its preface+SETTINGS
  - Streams are opened
- If the client has no prior knowledge, HTTP/2 has to be negotiated

# HTTP/2 Negotiation

- TLS-based Negotiation (HTTPS)
  - HTTP/2 can be negotiated via the TLS **Application Layer Protocol Negotiation (ALPN)**
    - client requests one of a list of protocols (including HTTP/2)
    - server selects one and informs the client in the response
    - if HTTP/2 has been selected, HTTP/2 handshake occurs as in the previous case (preface + SETTINGS exchange)
- HTTP-based Negotiation (HTTP)
  - HTTP/2 can be negotiated via the **HTTP1.1 Upgrade Mechanism**

# HTTP/2 Upgrade

- The client starts an HTTP/1.1 connection with upgrade header:

GET / HTTP/1.1

Host: server.example.com

Connection: Upgrade, HTTP2-Settings

Upgrade: h2c

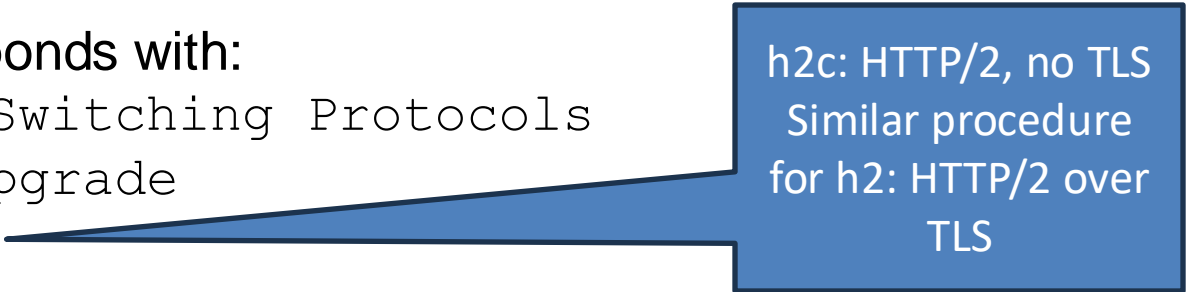
HTTP2-Settings: <base64 encoded HTTP/2 SETTINGS>

- The server responds with:

HTTP/1.1 101 Switching Protocols

Connection: Upgrade

Upgrade: h2c



h2c: HTTP/2, no TLS  
Similar procedure  
for h2: HTTP/2 over  
TLS

- The conversation continues as in the previous cases (preface + SETTINGS exchange)



# Error Management

- In addition to the HTTP/1 error reporting mechanisms (Response-level Error Reporting), HTTP/2 offers other error reporting mechanisms
  - Stream-level Errors occurring in the management of single streams
  - Connection-level Errors occurring in the management of the TCP connection
- The result is a more precise report, especially about non-completed requests

# Stream-level and Connection-level Error Management

- Two types of errors:
  - connection-level error (the TCP connection is compromised)
  - stream-level error (only one stream is compromised)
- Error codes:
  - NO\_ERROR (0x0)
  - PROTOCOL\_ERROR (0x1)                      generic protocol error
  - INTERNAL\_ERROR (0x2)
  - FLOW\_CONTROL\_ERROR (0x3)              flow control protocol violation
  - SETTINGS\_TIMEOUT (0x4)
  - STREAM\_CLOSED (0x5)
  - FRAME\_SIZE\_ERROR (0x6)
  - ...

# Other ways to learn about status of Non-completed Requests

- The GOAWAY frame indicates the highest stream number that might have been processed.
  - Requests on streams with higher numbers are therefore guaranteed to be safe to retry.
- The REFUSED\_STREAM error code can be included in a RST\_STREAM frame to indicate that the stream is being closed prior to any processing having occurred.
  - Any request that was sent on the reset stream can be safely retried.

# gRPC

- RPC mechanism that works on top of HTTP/2
  - exploits performance-oriented HTTP/2 features
    - binary messages, header compression, etc.
  - uses HTTP/2 as transport (like SOAP)
  - typical data representation: protocol buffers
- multi-language
- at most once semantics
- security management (e.g. via TLS)
- multiple ways of operation
  - synchronous, asynchronous, stream-oriented



# Protocol Buffers

- System-independent data representation and IDL
  - abstract language
  - binary representation
- Bindings for all major programming languages

# Protocol Buffers IDL v3

- C-like language (termination, blocks, comments)
- Example

```
syntax = "proto3";

service HelloService {
    rpc SayHello (HelloRequest) returns (HelloResponse);
}

message HelloRequest {
    string greeting = 1;
}

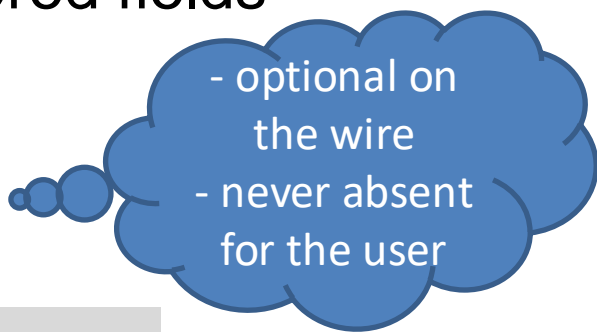
message HelloResponse {
    string reply = 1;
}
```

# Scalar Data Types

types	description
float, double	single/double precision floating point
int32, int64	variable-length encoding integer, inefficient for negative values
uint32, uint64	variable-length encoding unsigned integer
sint32, sint64	variable-length encoding signed integer, more efficient for negative values
fixed32, fixed64	fixed-length integer
sfixed32,sfixed64	signed fixed-length integer
bool	boolean
bytes	byte sequence (no longer than $2^{32}$ )
string	UTF-8 encoded or 7-bit ASCII string (no longer than $2^{32}$ )

# Message Types

- User-defined named data structures (a message to be used for RPC or part of it)
- Collection of typed, named, and numbered fields
  - each field can be
    - optional singular (default): 0 or 1 repetitions
    - optional repeated: 0 or more



- optional on the wire  
- never absent for the user

```
message Book {  
    string isbn = 1;  
    string title = 2;  
    repeated string authors = 3;  
}
```

- if a field is absent, its default value will be returned when the message will be read



# Enum Types

- Integers with variable-length encoding that can take a limited, well-defined set of named values

first value must be 0  
(it is the default)

```
enum ItemType {  
    ARTICLE = 0;  
    BOOK = 1;  
}  
  
message BiblioItem {  
    ItemType type = 1;  
    string title = 2;  
    repeated string authors = 3;  
}
```

# Default Values

Type	Default Value
numeric types	0
bool	false
bytes	the empty sequence of bytes
string	the empty string
enum types	the first value (0)
message types	unset (language dependent)

# Map Types

- Shortcut syntax for associative maps

```
message Biblio {  
    map<string, BiblioItem> items = 1;  
    string description = 2;  
}
```

- map fields cannot be repeated
- entry ordering is not relevant (it is not preserved on the wire)
- field `map<key_type,value_type> map_field = N;` is equivalent to

```
message MapFieldEntry {  
    key_type key = 1;  
    value_type value = 2;  
}  
  
repeated MapFieldEntry map_field = N;
```

# Oneof

- A set of fields such that at most one of them can be set to a nondefault value
  - the fields in the set share space on the wire

```
message BiblioItem {  
    string title = 1;  
    repeated string authors = 2;  
    oneof extension {  
        Article article = 3;  
        Book book = 4;  
    }  
}
```

- setting one field in the set will automatically unset all the others
- oneof fields cannot be **repeated** (and cannot be **maps**)

# Managing Large Specifications

- Import feature

```
import "biblio.proto" ;
```

- Packages

- A proto file can have a package declaration

```
package it.polito.dsp.biblio;  
message Biblio { ... }
```

- Name clashes can be avoided by using package declarations and package specifiers

```
import "biblio.proto";  
message Topic {  
    it.polito.dsp.biblio.Biblio biblio = 1;  
    ...  
}
```

# Nesting

- Message and Enum type declarations can be nested inside other Message types
- Nested declarations can be referred by the dot notation

```
message BiblioItem {  
    message Article {...}  
    message Book {...}  
  
    string title = 1  
    repeated string authors = 2;  
    oneof extension {  
        Article article = 3;  
        Book book = 4;  
    }  
}  
message Articles {  
    BiblioItem.Article article;  
    ...  
}
```

# RPC Types Supported by gRPC

- Classical Synchronous RPC
- Two-Way Asynchronous RPC
- Streaming mode (independently for request and response)

# Defining RPC Interfaces in Proto Files

- An RPC interface can be defined by a **service** block
  - each RPC procedure has its **rpc** declaration
    - Request message type
    - Response message type
    - **stream** qualifier can precede request or response type

```
service UpdateService {  
  rpc update(UpdRequest) returns (UpdResponse);  
  rpc updateCS(stream UpdRequest) returns (UpdResponse);  
  rpc updateSS(UpdRequest) returns (stream UpdResponse);  
  rpc updateDS(stream UpdRequest) returns (stream UpdResponse);  
}
```



# Example

- Interface Definition for the Bank Account Example:
  - addDeposit and addWithdrawal operations