

Synchronization Algorithms

© Riccardo Sisto, Politecnico di Torino

Reference for study:

Van Steen, Tanenbaum, "Distributed Systems", chapter 6

Synchronization

- Processes in a DS are asynchronous
- Sometimes, we need to synchronize them
 - a process must wait until another process completes another operation
 - some ordering of events must be enforced
 - some timing requirements must be satisfied
- Two possible approaches for decentralized systems:
 - synchronize physical clocks
 - logical clocks

Physical Clock Synchronization

- Different possibilities exist
 - UTC Receivers
 - synchronization can be very accurate (up to 50ns precision) but expensive
 - NTP
 - cheaper solution but not so accurate

	NTP accuracy
LAN	<1ms
public internet	10-50ms
public internet under congestion	>100ms

Logical Clocks


- For synchronization purposes, agreement on **ordering of events** is often sufficient
- Logical clocks are distributed algorithms that can be used to achieve this sort of agreement

The Happens-before Relation in a DS

- Definition:


$x \rightarrow y$ event x happens before event y

- Baseline: in a distributed system based on message passing, $x \rightarrow y$ can be observed in these cases:
 - x happens before y in the same process



The Happens-before Relation in a DS

- Definition: the Happens-before relation
 $\mathbf{x} \rightarrow \mathbf{y}$ event \mathbf{x} happens before event \mathbf{y}
- Baseline: in a distributed system based on message passing, $\mathbf{x} \rightarrow \mathbf{y}$ can be observed in these cases:
 - \mathbf{x} happens before \mathbf{y} in the same process
 - \mathbf{x} and \mathbf{y} are the events of sending and receiving the same message (by different processes)
 - transitive property holds



Happens-before and Causality

- The happens-before relation also captures the possibility of **causal relation** (causality) between two events

– if $x \rightarrow y$ then x and y **may be** causally related

in cause-effect relation (x cause, y effect)
OR y depends on/is influenced by x


– if $x \not\rightarrow y$
and $y \not\rightarrow x$ then x and y **are not** causally related
(we say they are **concurrent**)

Lamport Clocks

- Goal: assign timestamps (i.e., clock values) to events in such a way that if $\mathbf{x} \rightarrow \mathbf{y}$ then $C(\mathbf{x}) < C(\mathbf{y})$
- Algorithm:
 - each process P_i keeps a local time counter C_i
 - P_i timestamps each local event \mathbf{x} with the current C_i value $C_i(\mathbf{x})$, after having incremented C_i
 - P_i timestamps each message it sends with the timestamp $C_i(\mathbf{s})$ assigned to the send event \mathbf{s}
 - when P_i receives a message with timestamp C_r ,
 - if $C_r > C_i$, P_i sets $C_i = C_r$ (clock adjustment) and then it timestamps the receive event as usual

equivalent to
 $C_i = \max(C_i, C_r)$


Example



Discussion

- Some events (in different processes) are assigned the same timestamp
 - If we want to avoid it, we can include the process id into the timestamp: $C(\mathbf{x}) = \langle C_i(\mathbf{x}), i \rangle$ if event \mathbf{x} occurs in P_i
 - assumption: process ids are unique
 - $\langle C_i(\mathbf{x}), i \rangle < \langle C_j(\mathbf{y}), j \rangle$ iff $C_i(\mathbf{x}) < C_j(\mathbf{y})$ or $C_i(\mathbf{x}) = C_j(\mathbf{y})$ and $i < j$
- $\mathbf{x} \rightarrow \mathbf{y} \Rightarrow C(\mathbf{x}) < C(\mathbf{y})$ but **not** vice versa
- The algorithm can be implemented by
 - a timestamping function that increments and returns the counter
 - hooks in the send/receive functions that timestamp send/receive events and messages and adjust the counter

Example with Total Ordering



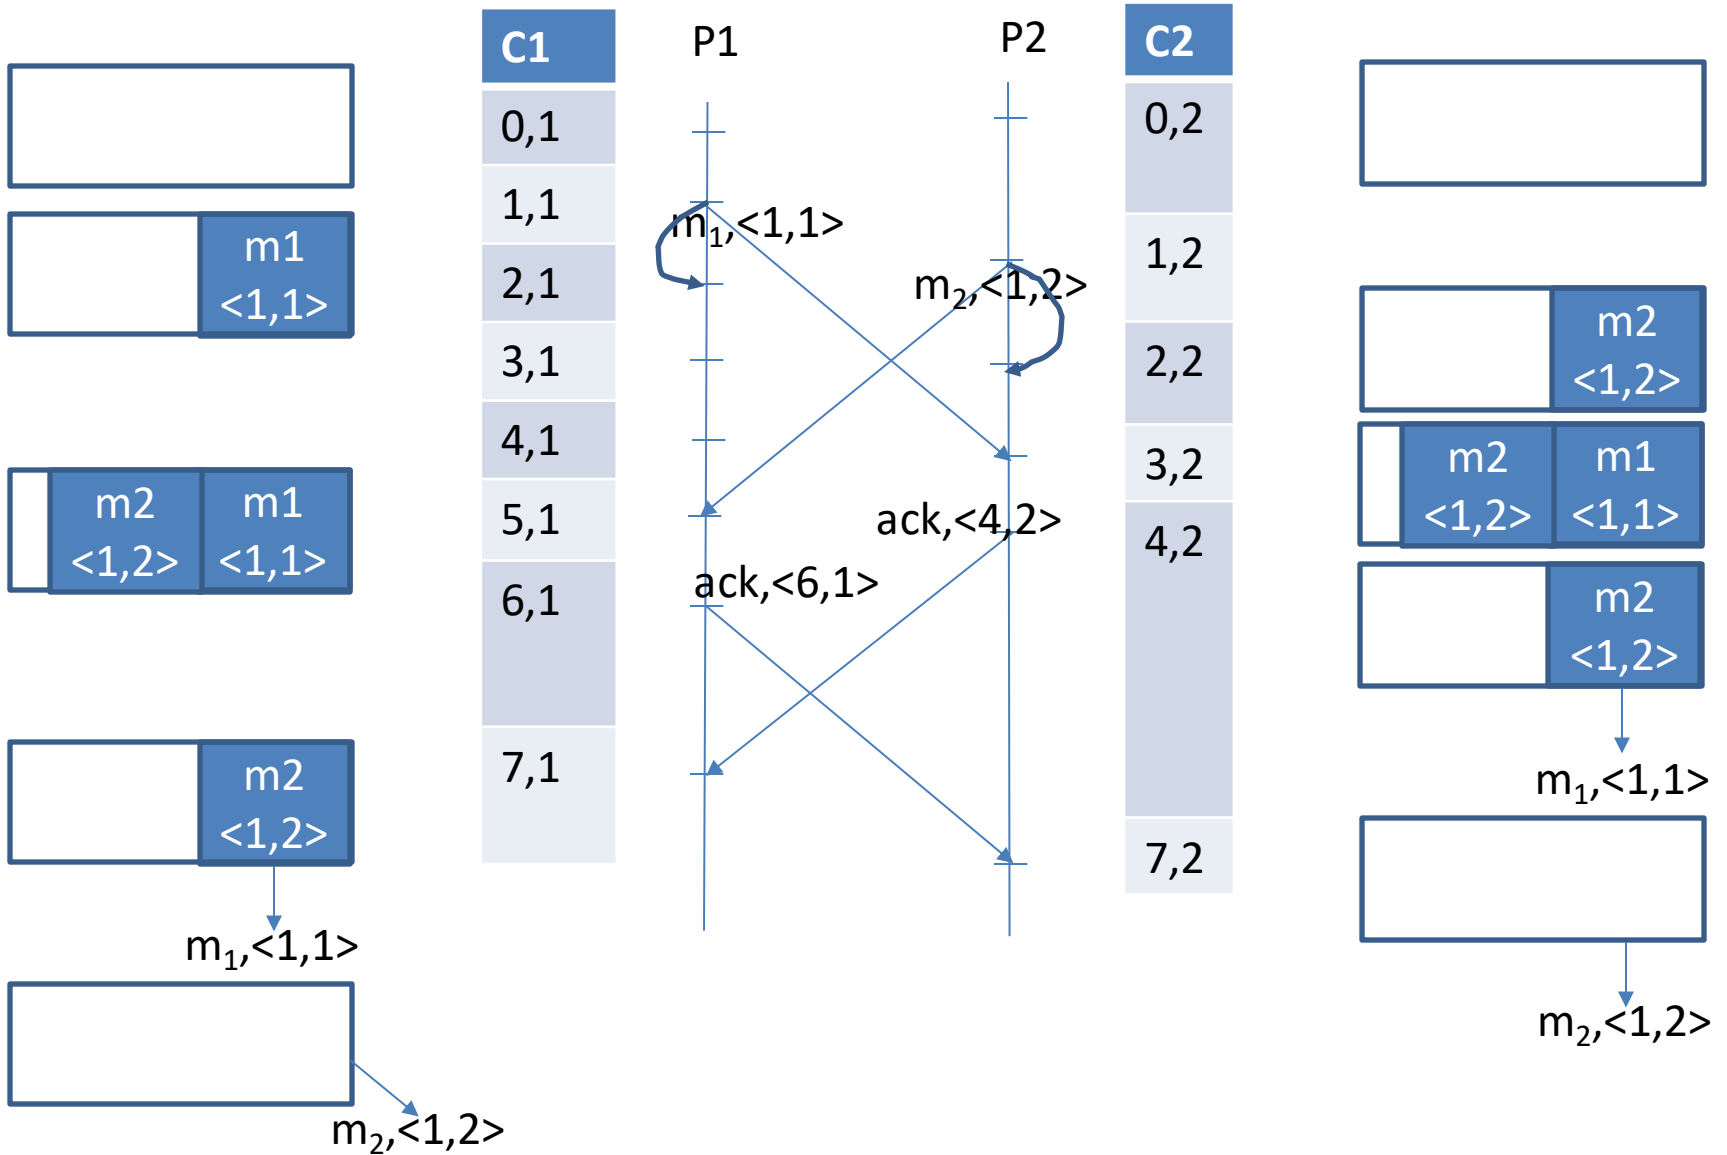
Application: Total-ordered Multicast

- Problem statement:
 - a set of processes that send multicast messages
 - all multicast messages must be delivered in the same order to each receiver (total-ordered multicast)
- Example
 - 2 or more identical replicated servers
 - replication goals: fault-tolerance and latency reduction
 - service requests can be sent to any server in the group
 - the receiving server multicasts the request to all servers
 - operations must be executed in each replica in the same order

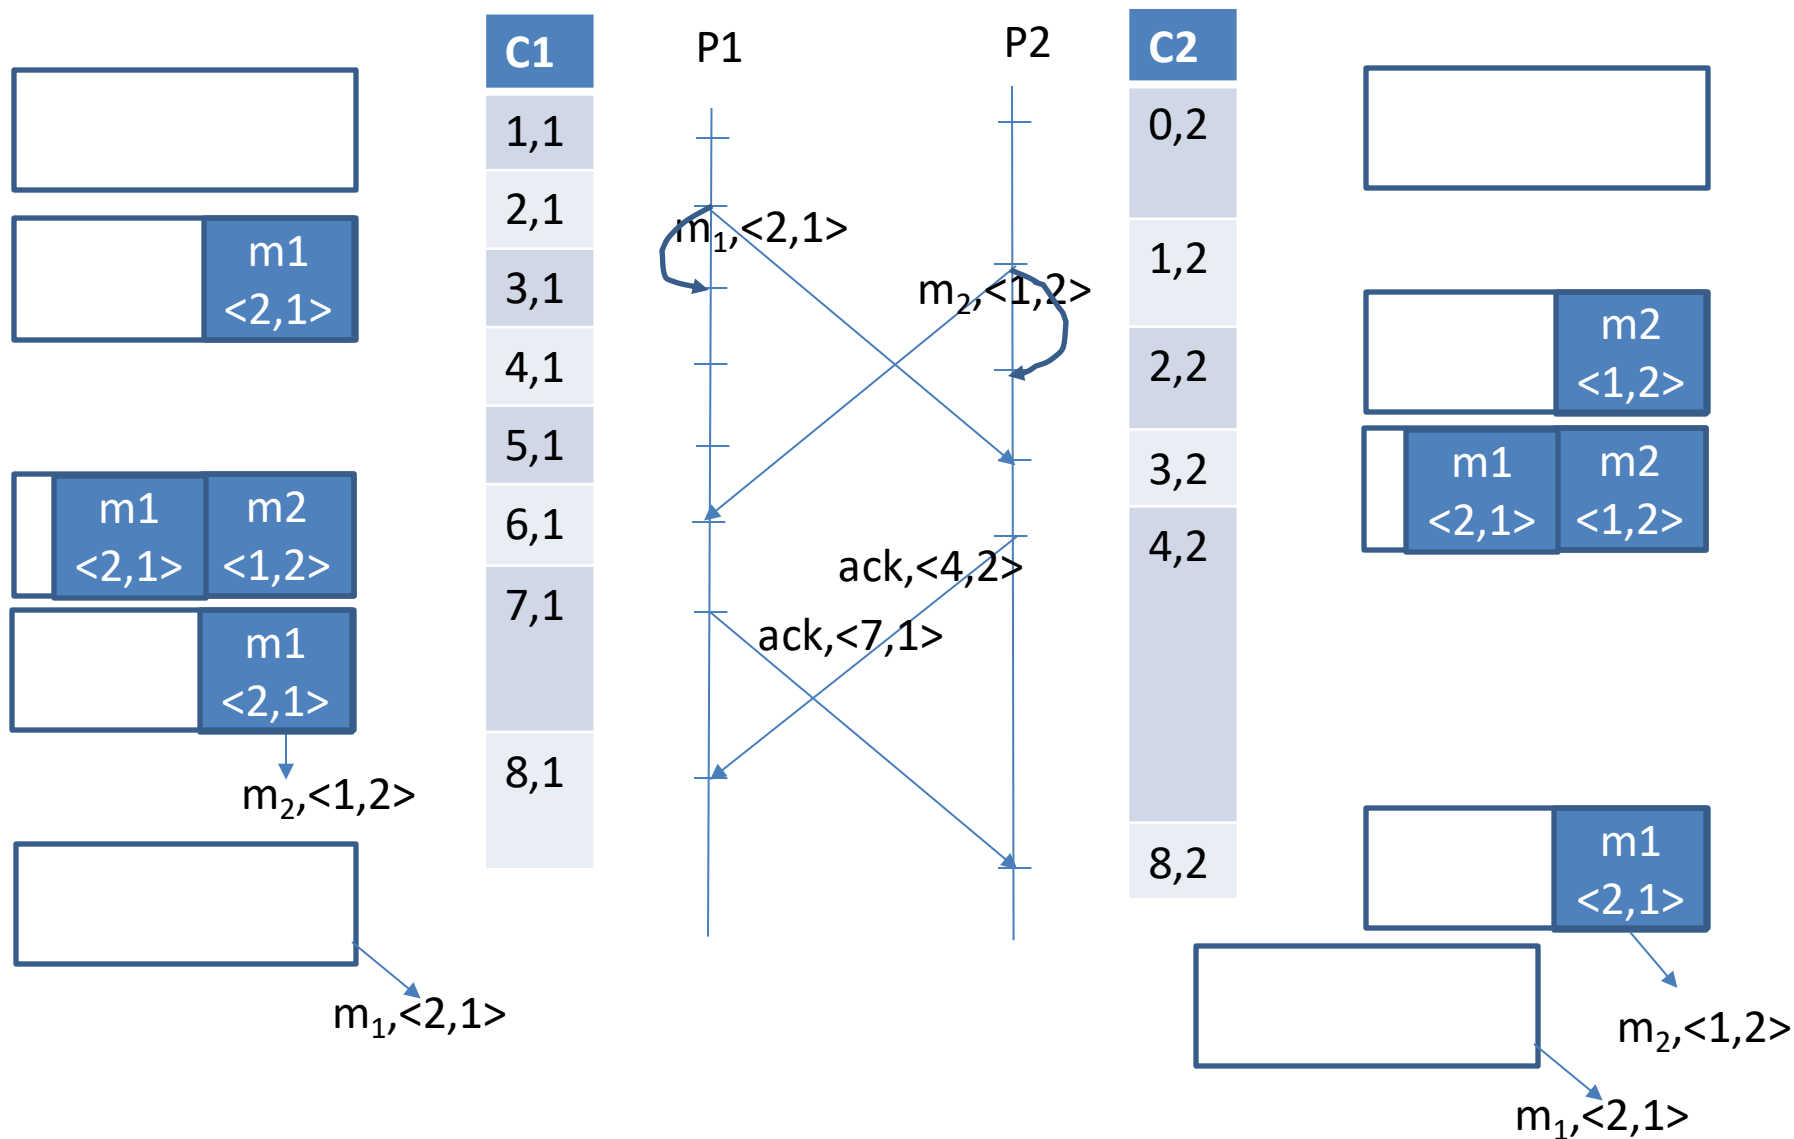
Total-ordered Multicast Algorithm

- Lamport clocks are maintained in each process.
 - => Each multicast message is timestamped by the sender with the send event timestamp
- Received messages are queued, ordered by timestamp
- Each receiver acknowledges message reception to the other receivers
- When acknowledgments for the message at the head of the queue have been received from all other receivers, that message is dequeued and handed to the application

Example



Example: Other Scenario



Total-ordered Multicast: Discussion

- The algorithm has been shown to deliver messages to all destination processes in the same order under the following assumptions:
 - no messages are lost
 - messages from the same sender are received in the same order they were sent
- In case the messages have the meaning of operations to be executed on replicated data in the same order, the algorithm implements state machine replication

Vector Clocks


- Lamport clocks provide a **total ordering** of events, but this ordering does not necessarily imply *causality*

$$C(x) < C(y) \quad ? \quad x \rightarrow y$$

- Example:

$$C(e) < C(i)$$

but not $e \rightarrow i$




- Vector clocks are a way to provide a **partial ordering** of events that captures causality, i.e.

$$C(x) < C(y) \iff x \rightarrow y$$


Vector Clocks Algorithm

- Each process P_i keeps a local **vector** of time counters $VC_i[]$
 - $VC_i[i]$ is the local event counter (i.e. local time at P_i , as with Lamport)
 - $VC_i[j]$ is P_i 's knowledge of the local time at P_j
- P_i timestamps each local event x with the current VC_i value $ts(x)$, after having incremented $VC_i[i]$
- P_i timestamps each message it sends with the timestamp $ts(s)$ assigned to the send event s
- when P_i receives a message with timestamp tsr ,
 - P_i sets $VC_i[k] = \max(VC_i[k], tsr[k])$ for each k (vector clock adjustment) and then it timestamps the receive event as usual

Example: How does this scenario change with Vector Clocks?



Example with Vector Clocks



Vector Clock Application Example

- Causal Ordered Multicast
 - each multicast message must be delivered after the delivery of all other causally related messages that were sent before
 - non causally related messages that were sent before may be received after
- Use Example
 - in a replicated chat, we may accept that the ordering of posts is different in each replica provided each reply to a post p does never come before p