

# **Coordination Algorithms**

© Riccardo Sisto, Politecnico di Torino

Reference for study:


Van Steen, Tanenbaum, "Distributed Systems", chapter 6

# Mutual Exclusion

- Problem Statement
  - Guarantee mutual exclusive access to shared resources by multiple processes in a DS
- Different types of algorithms
  - Token-based
  - Permission-based
    - centralized
    - decentralized

# Token Ring Mutual Exclusion

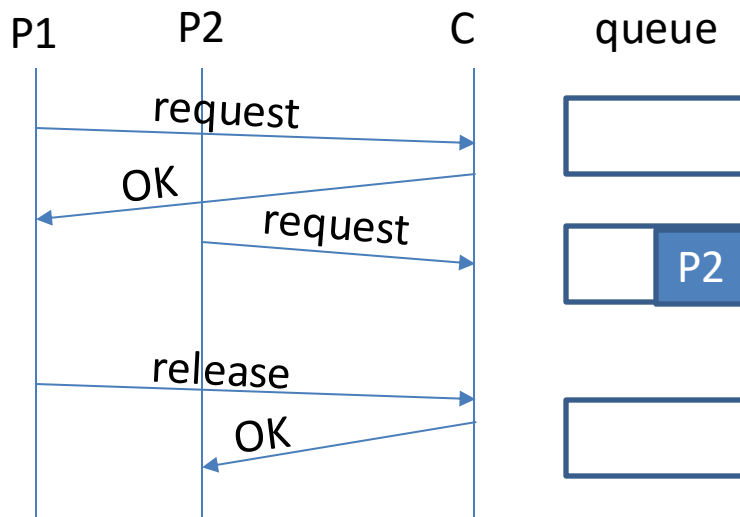
- Processes organized in a ring overlay
- 1 token in the system, continually circulating on the ring



- A process  $P$  willing to access the resource waits for the token. When the token arrives at  $P$ :
  - $P$  starts accessing the resource and keeps the token until its access is finished
  - Then,  $P$  passes the token to the next process in the ring

# Centralized Mutual Exclusion

- Central manager (C) of shared resource(s)
- A process P willing to access a resource sends request to C and waits for permission response
- C delays permissions while the resource is engaged



# Decentralized Mutual Exclusion

- Based on Lamport clocks (totally ordered multicast)
- Requester sends request to all processes (including itself) and waits for permission from every process
  - responses to requests received while accessing the resource are delayed
  - in case of conflict (receiver also wants to access the resource), the request with the lower timestamp wins

# Performance Comparison

Algorithm		#Messages/access	delay before entry (# messages)
Permission based	Centralized	3	2
	Decentralized	2 (N-1)	2 (N-1)
Token based	Token ring	1...	0 ... N-1

N: number of processes

# Election






- Problem Statement:
  - Elect a process (e.g., a coordinator) in a group of processes
- Assumptions:
  - Each process has a unique identifier  $\text{id}(P)$  (it can be obtained by means of a naming system)
  - Each process knows all the other processes in the group
  - Processes in the group can be **up and running** or **down**, but channels are reliable
- Requirements
  - The algorithm must elect the up process having the highest id
  - At the end of the algorithm, all processes agree about who is the elected process

# Election: The Bully Algorithm

- Let  $\text{id}(P_k)=k$
- The algorithm starts when a process  $P_k$  detects the coordinator is missing and decides to hold an election:
  - $P_k$  sends ELECTION message to  $P_{k+1}, P_{k+2}, \dots, P_{N-1}$
  - If no one responds,  $P_k$  wins the election, else  $P_k$  gives up
- Whenever  $P_i$  receives an ELECTION,
  - $P_i$  responds with OK (means it is alive)
  - $P_i$  holds an election (if it was not already holding one)
- Eventually all but one processes give up and one wins
- The winner finally informs the other processes



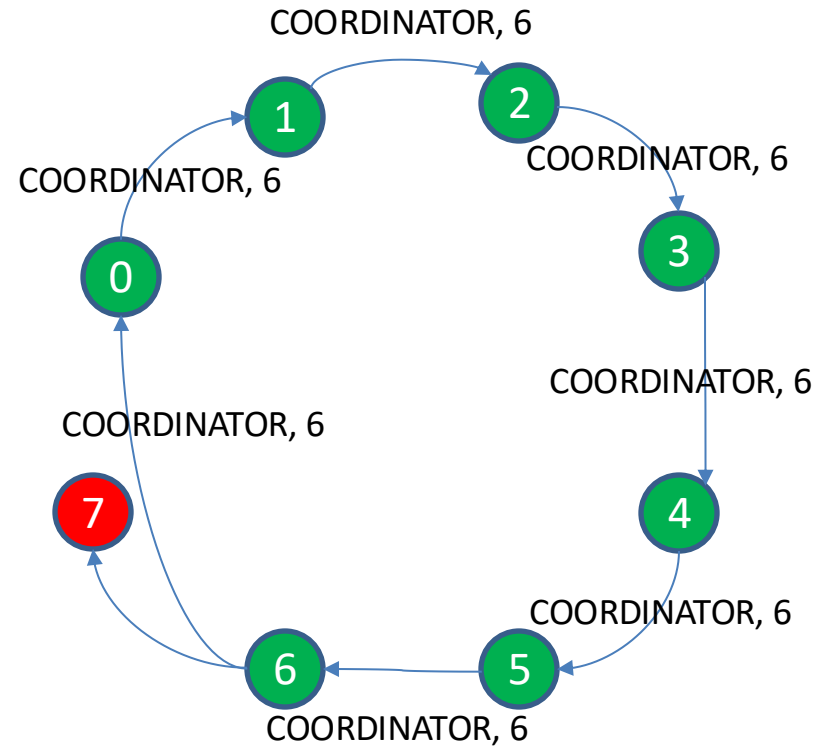
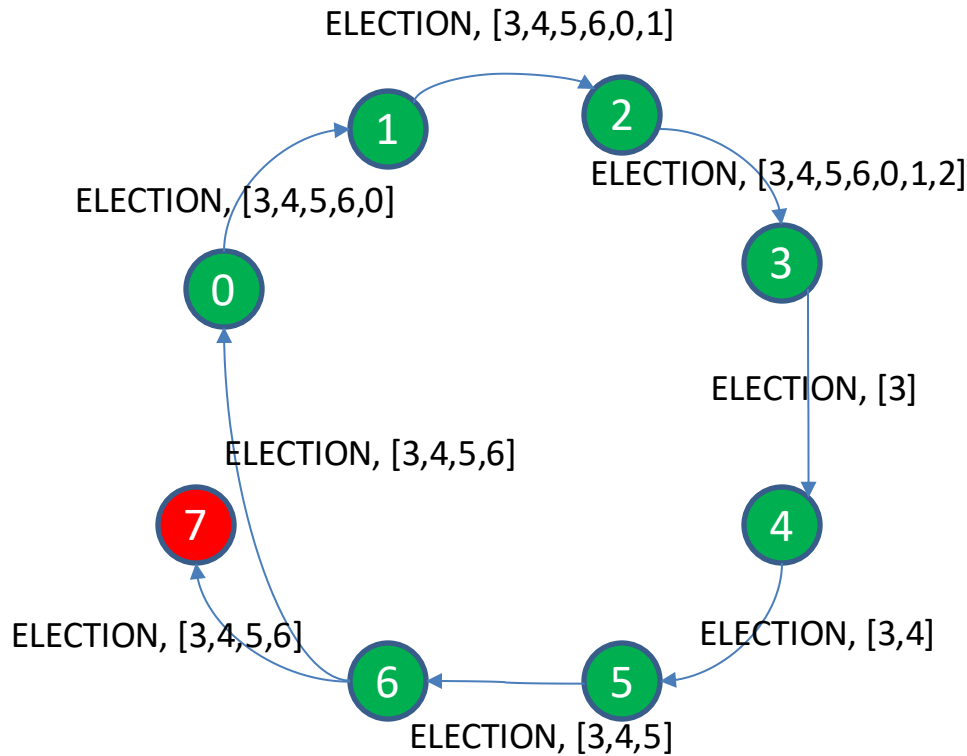
# Election: The Bully Algorithm



# Election: Ring Algorithm

- Processes are ordered (e.g., by id) in a logical ring
- The algorithm starts when a process  $P_k$  detects the coordinator is missing and decides to hold an election:
  - $P_k$  sends ELECTION message to its successor in the ring
  - if no response is received,  $P_k$  sends the message to the next successor in the ring and so on until one responds
- Each ELECTION message carries the list of senders
- When eventually the message gets back to the starter  $P_k$ 
  - $P_k$  stops the circulation of the message
  - $P_k$  computes the winner and circulates a COORDINATOR msg on the ring (in the same way) to inform about the winner

# Election: Ring Algorithm



# Consensus

- A more general coordination problem
- Problem statement
  - Let  $n$  processes, each one proposing an input value, agree on the same output value
- Properties
  - Mutual exclusion, leader election: special cases of consensus
  - But consensus can be solved by using leader election or mutual exclusion algorithms