

# TCP/IP Sockets

© Riccardo Sisto, Politecnico di Torino


References for study: K.L. Calvert, M.J. Donahoo  
**TCP/IP Sockets in Java, 2nd Edition**, O'Reilly 2011

# The Socket API

- The de-facto standard API for accessing network services provided by Internet layers 4-3-2
- Born with **BSD Unix**, in the **C language**, but later implemented for all major OSs and languages
- The name comes from the main abstraction of the API: the **socket** (an endpoint of a layer 4-3-2 protocol)

# General Features of the Socket API

- Procedural API in the C language
- Main idea: extend the conventional Unix I/O model:



- One model for several protocol stacks
- Later implementations have partially diverged from these original features

# Sockets

- A **socket** is the abstraction of an inter-process communication channel endpoint or SAP
- Sockets “live” in **domains**, each one characterized by its **protocol family** and **address family**
- Communication between two different domains is impossible.
- Domain examples:

Domain	Protocol Family	Address Family
ARPA Internet	PF_INET	AF_INET
Internet with IPv6	PF_INET6	AF_INET6
ISO/OSI	PF_ISO	AF_ISO
Unix pipes	PF_UNIX	AF_UNIX

# Socket Features: the TYPE

- The TYPE of a socket specifies the *service type* accessible through the socket. The main supported types are:
  - **SOCK\_STREAM**  
**continuous bidirectional byte stream** (without delimiters), transmitted reliably (same byte order, no duplications) (**connection-oriented** service, offered by layer 4, e.g. TCP).
  - **SOCK\_DGRAM**  
**bidirectional message** (datagram) **delivery** without reliability guarantee (messages can be delivered out of order and can be duplicated) (**connectionless** service, offered by layer 4, e.g. UDP).
  - **SOCK\_RAW**  
direct access to the services provided by layer 2-3 protocols

# Socket Features: the PROTOCOL

- In each domain, and for each socket type, a protocol to be used with the socket can be selected. For example, in the PF\_INET domain the following choices are possible:

## **SOCK\_STREAM type**

- TCP Protocol (IPPROTO\_TCP)

## **SOCK\_DGRAM type**

- UDP Protocol (IPPROTO\_UDP)

## **SOCK\_RAW type**

- ICMP Protocol (IPPROTO\_ICMP)
- IP Protocol (IPPROTO\_RAW)

# Socket Features: the OPTIONS

- Options specify various other socket features

Examples:

- `SO_RCVBUF` (size of receiver buffer)
- `SO_SNDBUF` (size of transmitter buffer)
- `SO_LINGER` (if enabled, delays connection closing when undelivered data are present in the socket buffer)
- `SO_KEEPALIVE` (if enabled activates transmission of periodic keepalive messages: if the response to a keepalive message is not received, the connection is closed)

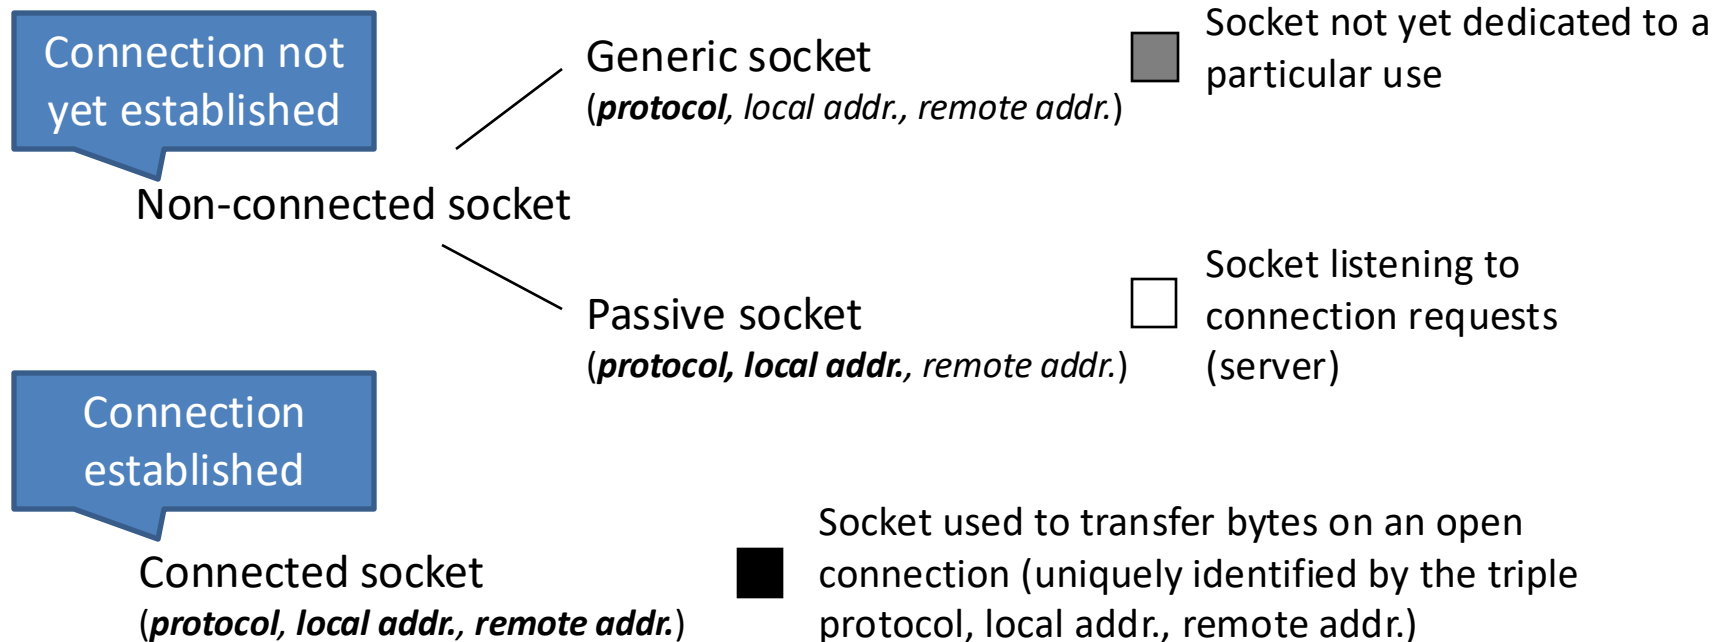
# Data Structure Associated with a Generic Socket

domain
type
protocol
local address
remote address
Options

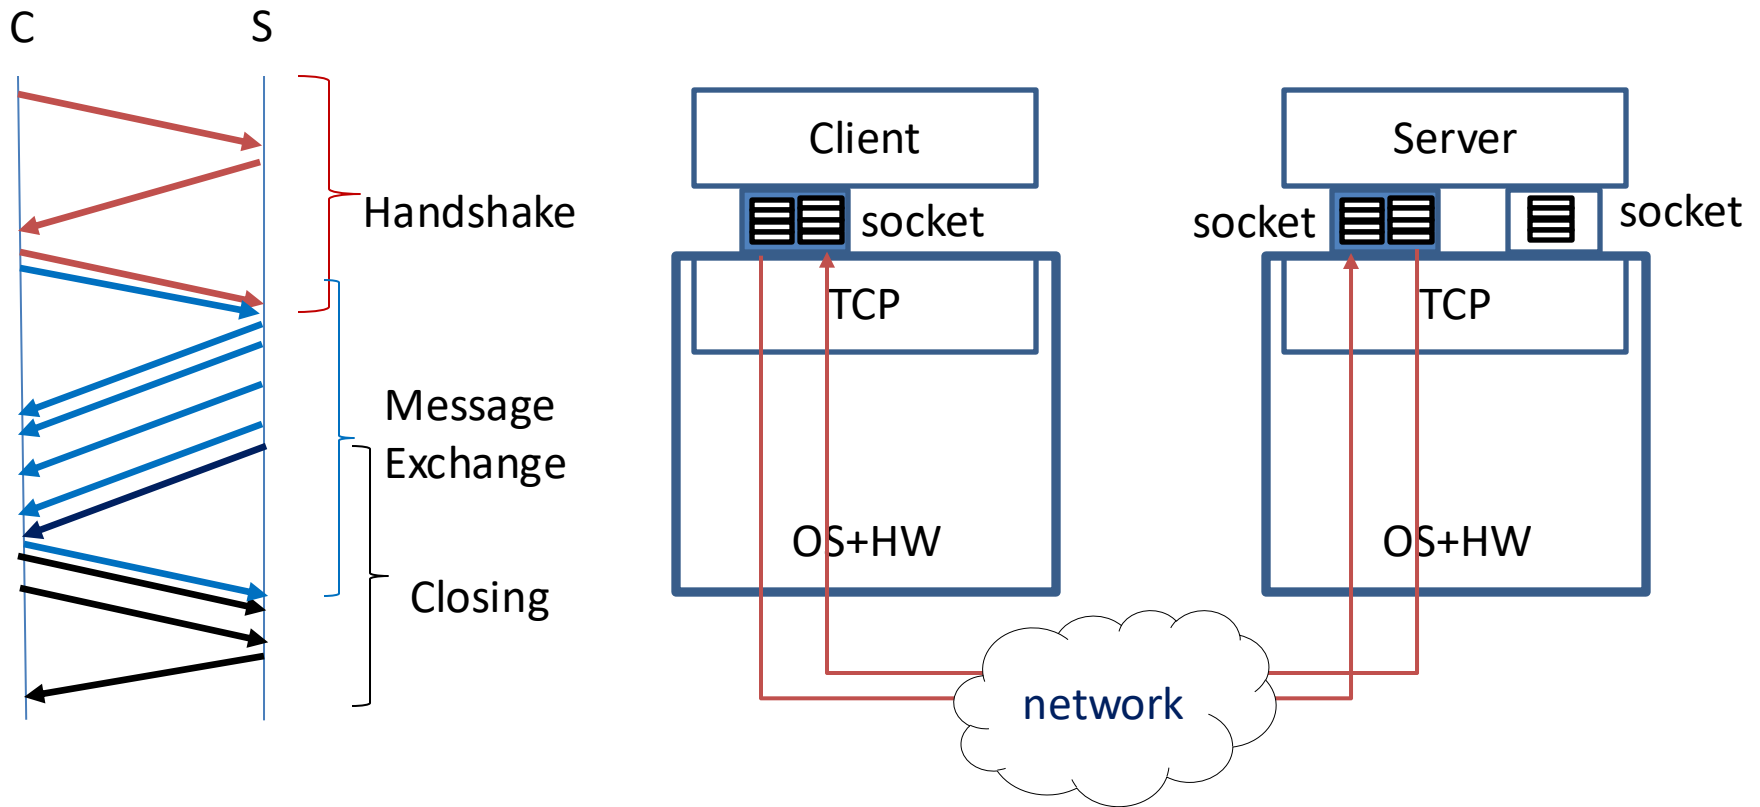


# Using TCP (STREAM) sockets

- Endpoints can take different states/forms, according to the role they are playing in communication:



# Connected and Passive Sockets




# Operations Offered by the Socket API for Accessing Layer 4 (TCP)

- Allocate local resources for communication
- Specify endpoints
- Open a connection (client-side)
- Wait for the establishment of a connection (server-side)
- Send/receive data on a connection (including urgent data)
- Be notified when data arrive
- Gracefully terminate a connection or abort a connection
- Respond to graceful termination requests and abort conditions
- Release resources when a connection terminates

# The Java Implementation of the Socket API

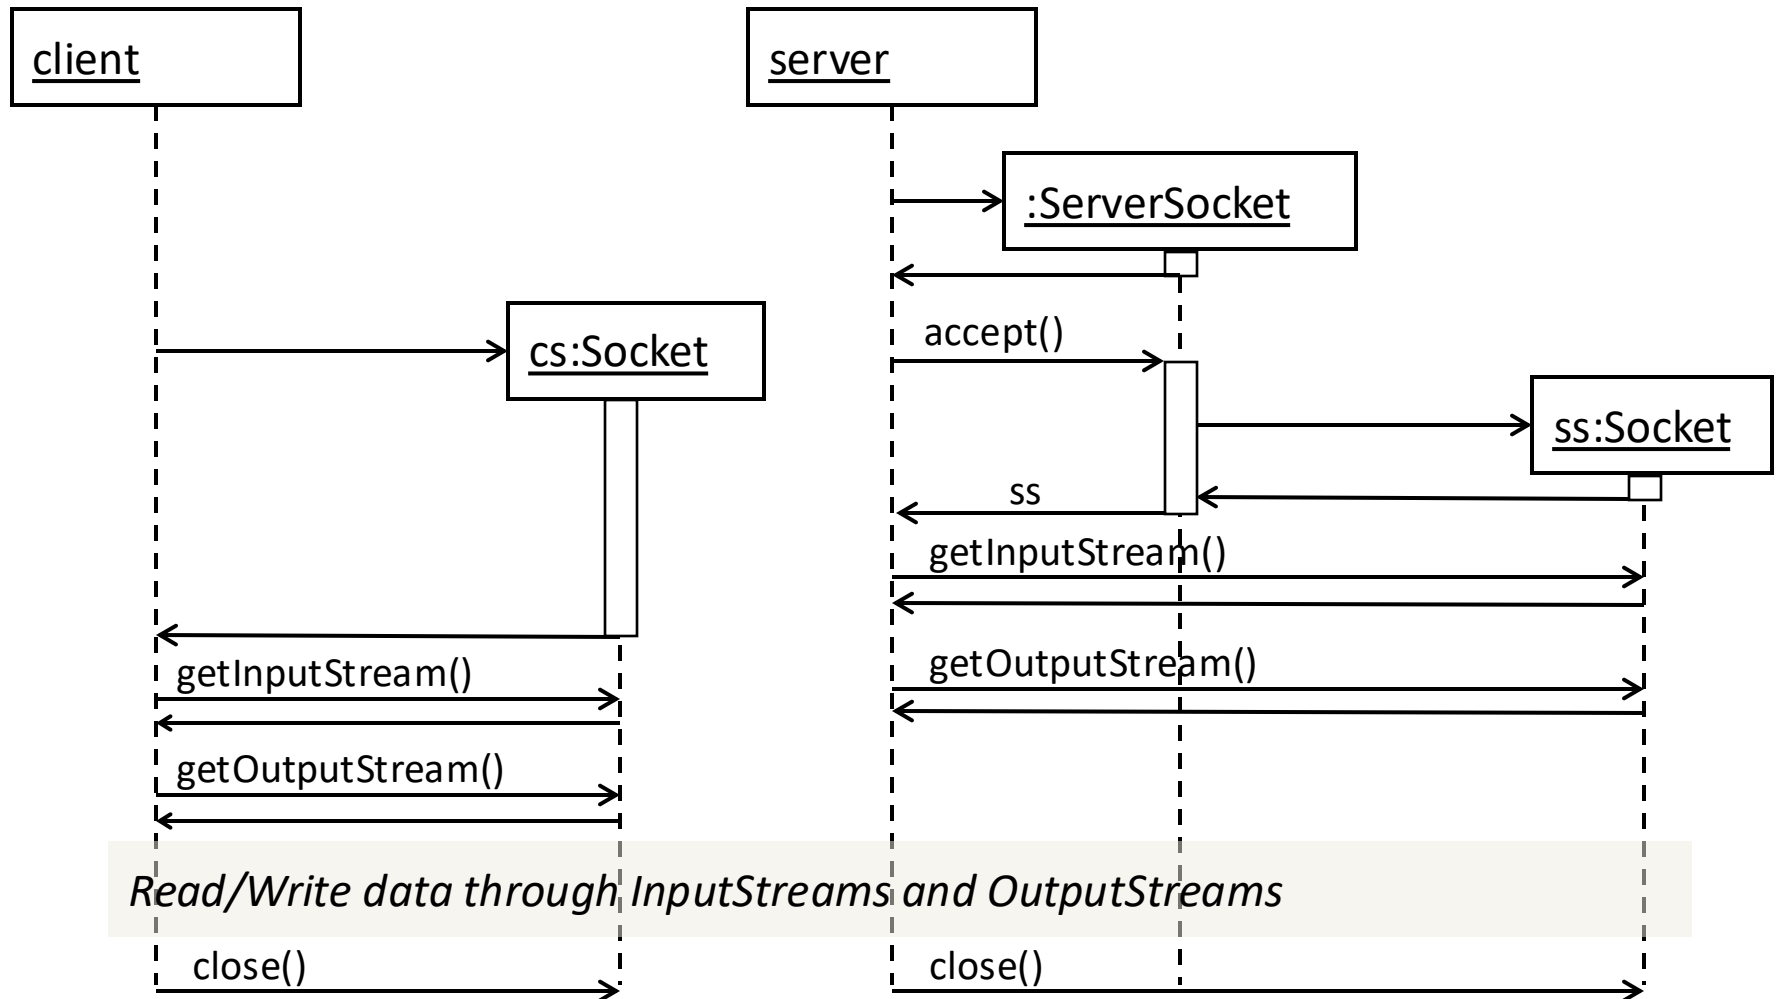
- The Java Socket API includes all the features of the original Socket API, with a special focus on the Internet domain (TCP/IP sockets)
- The core classes are in the **java.net** package



# The java.net Package

Class	Meaning
<b>NetworkInterface</b>	interface information (name, IP addresses, etc)
<b>InetAddress</b>	Internet address (with methods for DNS services)
<b>SocketAddress</b>	Generic socket address (no address family)
<b>InetSocketAddress</b>	Internet socket address
<b>Socket</b>	Connected STREAM socket
<b>ServerSocket</b>	Passive STREAM socket
<b>DatagramSocket</b>	DGRAM socket
<b>DatagramPacket</b>	Datagram

# Using Connected Sockets



# Establishing a Connection

- Server-side

All parameters are optional

- create ServerSocket

ServerSocket(int port, int backlog, InetAddress bindAddr)

- call accept() to wait for connection establishment

=> returns Socket when connection established

- Client-side

- create Socket

Socket(InetAddress address, int port, InetAddress localAddr, int localPort)

Socket(String host, int port, InetAddress localAddr, int localPort)

# Exceptions

java.io.IOException

|-- java.net.SocketException

| |-- BindException,

| |-- ConnectException,

| |-- NoRouteToHostException,

| |-- PortUnreachableException

|

|-- java.net.UnknownHostException

java.lang.SecurityException



# Closing a Connection

- **Socket.close()** terminates *both* sides of the connection
  - a FIN is sent by TCP after all buffered data have been delivered
  - a FIN from the peer will be automatically acknowledged by TCP
  - no further possibility to read from or write to the socket
- **Socket.shutdownOutput()** terminates the output side
  - a FIN is sent by TCP after all buffered data have been delivered
  - no further possibility to write to the socket
- **Socket.shutdownInput()** terminates the input side
  - undelivered data are discarded
  - a FIN from the peer will be automatically acknowledged by TCP
  - no further possibility to read from the connection

# How to Detect Closing

- A read operation on the socket input stream returns -1 if the incoming side of the connection has been closed by the peer
  - when this happens, the process may want to start the closing of the other side of the connection (after having sent any outstanding data)

# Example: Echo Protocol on TCP

- Implement the simple echo service provided by Unix systems (on standard port 7):
  - the server continuously echoes back the bytes received on the connection
  - the connection is closed by the client when done

=> [TcpEchoServer0.java](#)
- Implement a client for the service
  - the client operates with line buffering

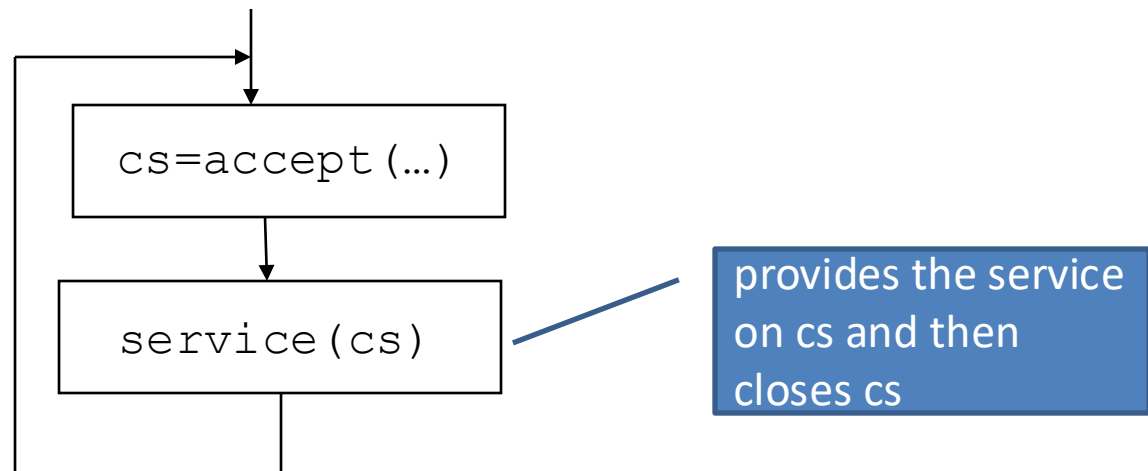
=> [TcpEchoClient0.java](#)

# Blocking Operations and Timeouts

- Some operations are blocking
  - accept
  - new Socket (connect)
  - read operations on InputStream (block until 1 byte available)
  - write operations on OutputStream (block until space available)
- When blocking is undesired
  - the operation can be called by a dedicated thread
- Timeouts for limiting read blocking can be set by setting the SO\_TIMEOUT socket option

# Sequential Servers

- The **sequential** model is the simplest one: the server serves requests sequentially in FIFO order:
  - only when a request has been completely served, the next one is processed
  - for TCP servers, the algorithm has the following structure:



# Problems of Sequential Servers

- Each request must wait until all previous requests have been served
- Waiting occurs even when the CPU is idle!
- When the request rate is high (time between requests and service time are comparable) it is likely that a request arrives while the server is busy.

⇒ requests may be discarded, or clients may timeout.


Possible solution: serve requests concurrently

# Concurrent Servers

- May serve more requests concurrently:
  - as a request arrives its service is started and proceeds concurrently with the other ones
- The probability that a request arrives when the server cannot accept it is reduced
- There are different ways to implement a concurrent server:
  - by assigning each request to a different thread
  - by simulating concurrency within a single thread

# TCP Concurrent Server with Thread Creation on Demand

- A Master Thread continuously accepts connections
- As a new connection is accepted a slave thread is created to serve it




[TcpEchoServer1.java](#)



# Solution Based on Pre-Created Thread Pool

- In order to reduce response time without giving up the advantages of concurrency, a pool of threads can be created at startup and allocated to requests when they arrive



[TcpEchoServer2.java](#)

# Java Executors

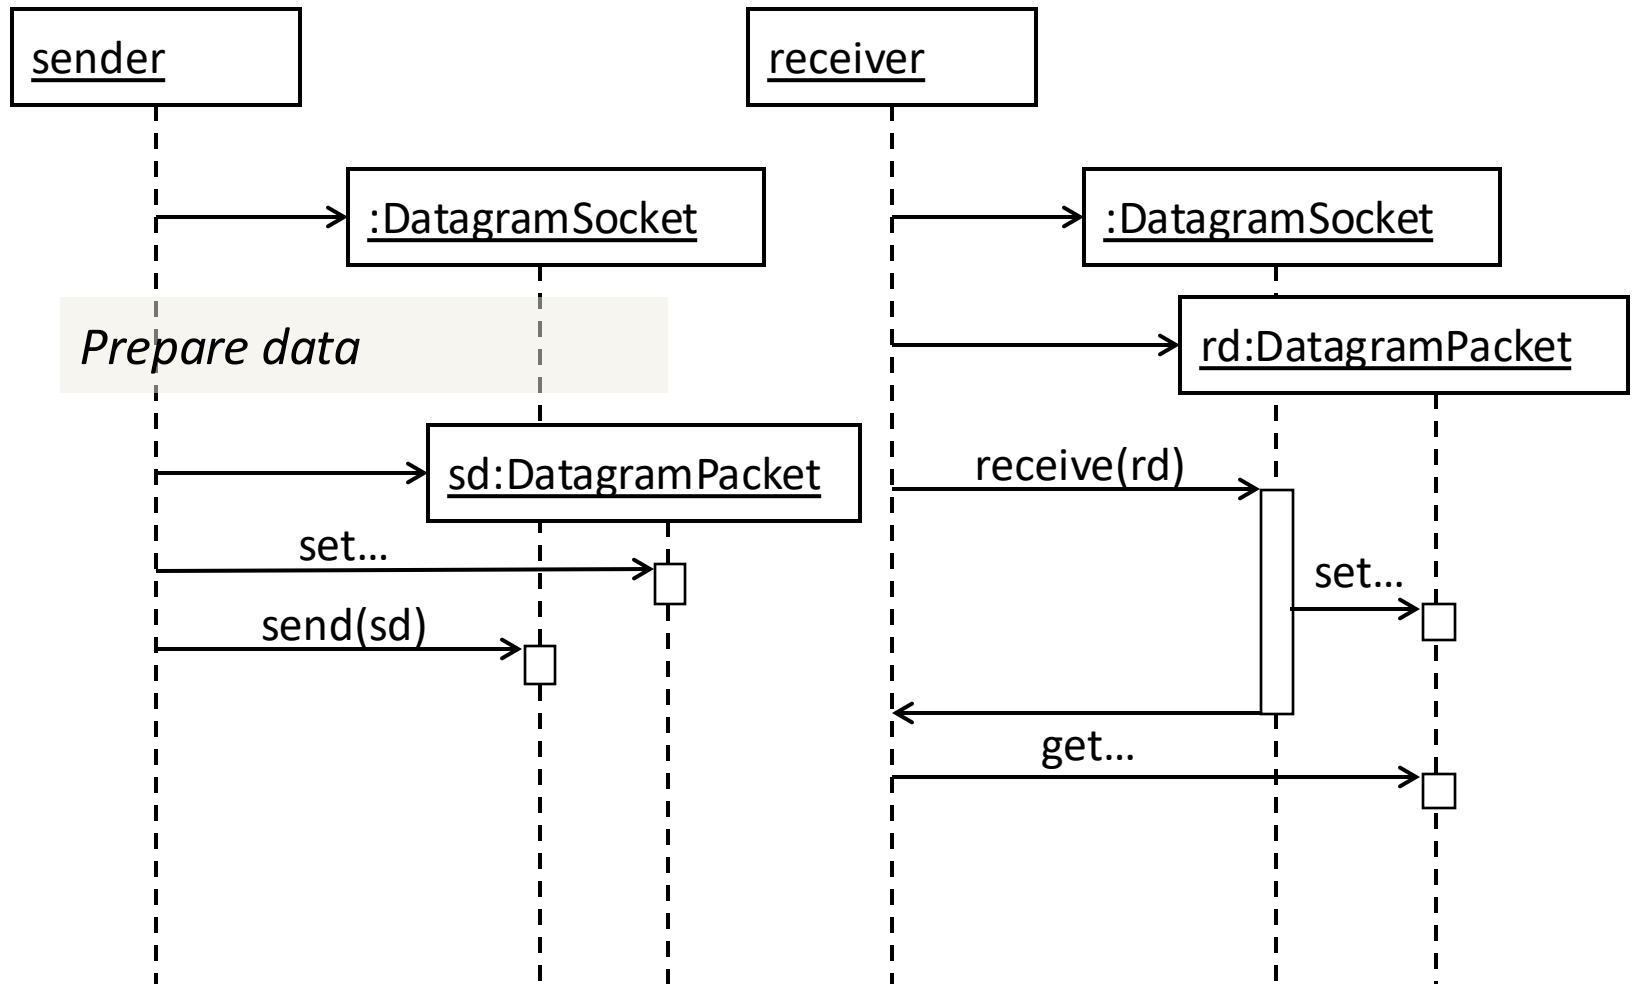
- Java supports thread pool management through Executors (in the `java.util.concurrent` library)
  - different executor types available
  - possibility to configure thread life-cycle (e.g. maximum/minimum number of threads), policies for re-launching failing threads,...

[TcpEchoServer3.java](#)

# Operations Offered by the Socket API for Access to Layer 4 (UDP)

- Allocate local resources for communication
- Specify endpoints
- Send/Receive a datagram
- Release allocated resources

# Using UDP Sockets in Java



# Using UDP Sockets in Java

- DatagramSocket

optional

DatagramSocket(int localPort, InetAddress localAddr)

- possibility to "connect" the socket to a remote address

connect(InetAddress remoteAddr, int remotePort)

- close() destroys socket and releases resources

- DatagramPacket

optional

DatagramPacket(byte[] buf, int length, InetAddress destAddr, int destPort)

# Example: Echo Protocol on UDP

- Implement the simple echo service provided by Unix systems (on standard port 7):
  - the server continuously echoes back the datagrams received  
=> [UdpEchoServer0.java](#)
- Implement a client for the service
  - the client operates with line buffering  
=> [UdpEchoClient0.java](#)

# Encoding/Decoding Data

- TCP and UDP provide transport of bytes
- Data encoding and decoding (marshalling/unmarshalling) is up to the programmer
  - possibility to use libraries for standard encoding formats (XML, JSON, Protocol buffers, ...)
  - in Java, other data encoding/decoding functions available in library classes (e.g., `DataInputStream`, `DataOutputStream`, `String`)
    - care necessary: e.g., integer endianness, charset