

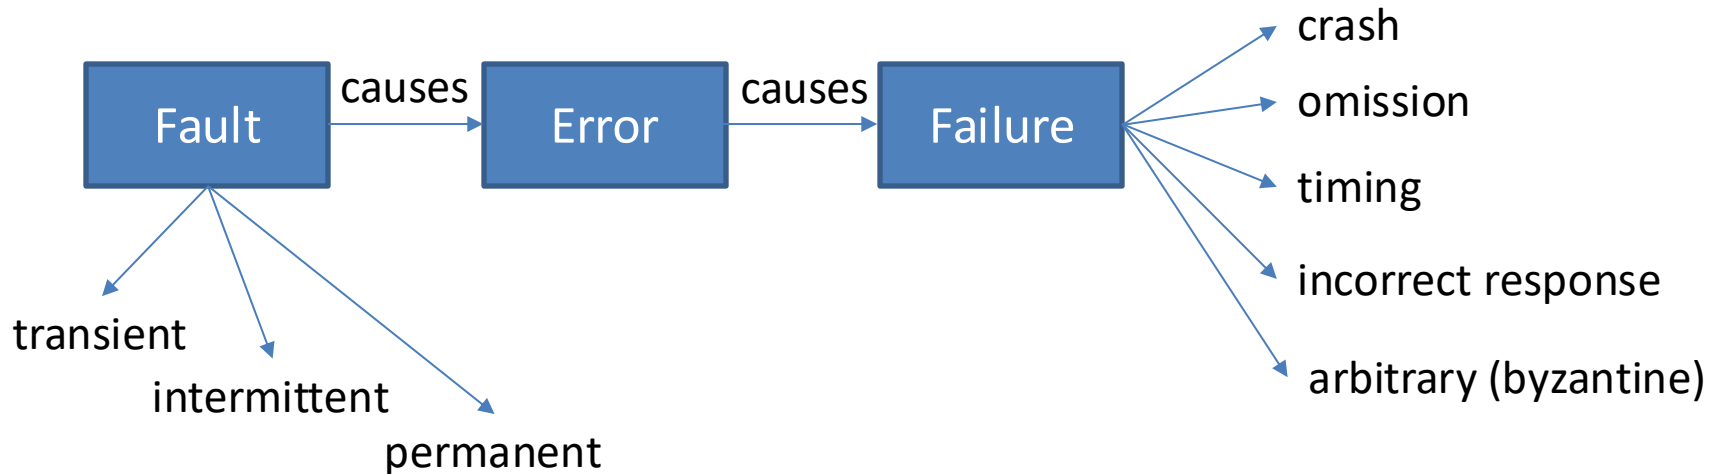
Fault Tolerance

© Riccardo Sisto, Politecnico di Torino

Reference for study:

Van Steen, Tanenbaum, "Distributed Systems", chapter 8

Main Fault Tolerance Concepts



- Dependability
 - Availability, Reliability, Safety, Security, Maintainability
 - Threatened by faults
 - Can be achieved by a combination of
 - fault **prevention**, fault **tolerance**, fault **removal**, fault **forecasting**

Fault Tolerance of Distributed Systems

- Ability to continue delivering service in the presence of partial failures, possibly with degradation of service, but without seriously affecting the overall performance
 - automatically **recover** from partial failures
 - failures need to be **detected**
 - recovery procedures must be implemented
 - **continue operation** in the presence of partial failures (while waiting for **repair**), rather than stopping

Detection, Recovery and Repair Techniques

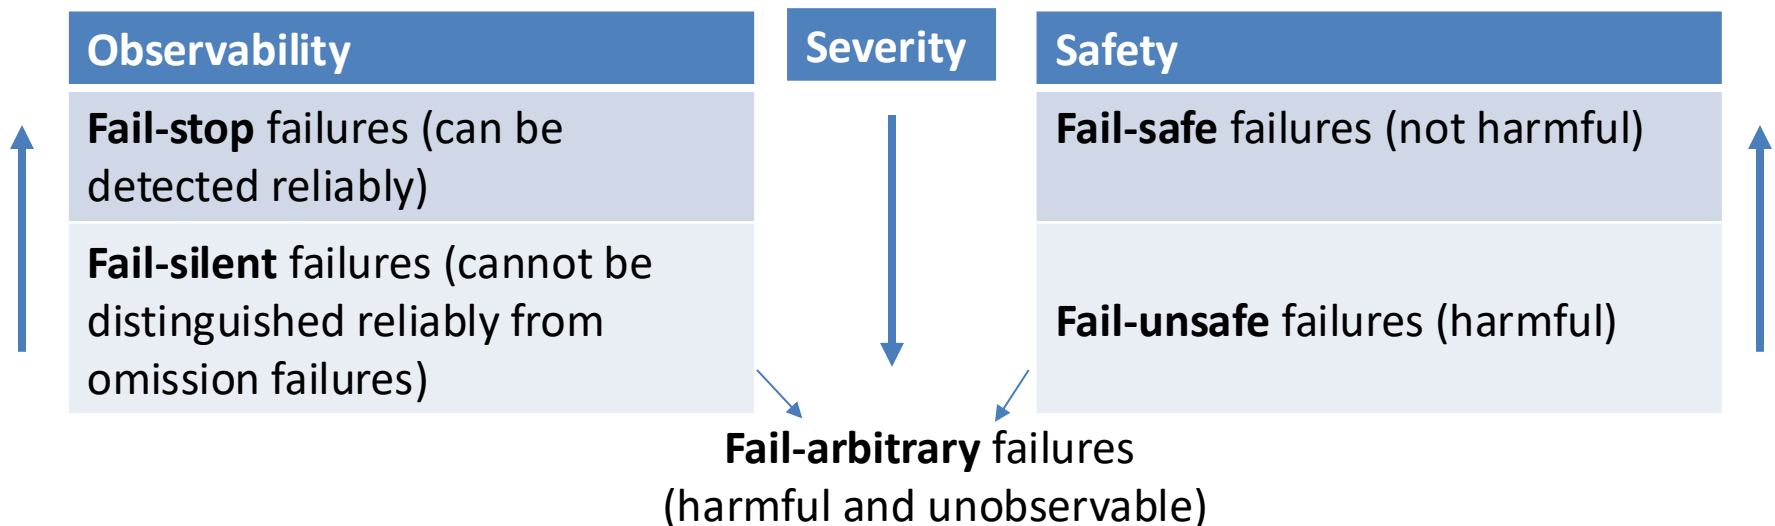
Failure Type	Detection Techniques	Recovery/Repair Techniques
crash	timeout	replacement, restart, redundancy
omission	timeout	redundancy (including retransmission)
timing	timeout	redundancy
incorrect response	input validation /comparison	redundancy
arbitrary	input validation /comparison	redundancy

Redundancy Techniques

- Time redundancy
 - retransmit messages
 - redo operations that fail (e.g., retry an aborted transaction)
- "Physical" Redundancy
 - adding more equipment, processes or data copies
 - The most common solution: process replication (process groups)

Detecting Process Crash

- DS are asynchronous with no delay upper bound
 - A process P can detect the crash failure of another process Q by timing out when no data arrive from Q, but detections may be **erroneous** and **delayed**
- Not all crash failures have the same observability/safety/severity:



Example: Detecting Crash of Process Connected via TCP

- The TCP layer itself can **know** that a connection has been closed or reset by the peer
 - FIN or RST segment received
- or **guess** that communication with the peer has been lost
 - Maximum number of retransmissions reached => Reset
- A process is informed about these conditions by the Socket API when reading or writing:

	Read	Write
connection closed	returns -1	successful
connection reset	throws SocketException	throws SocketException

Possible Crash Scenarios


- Process crash
 - OS detects crash, TCP layer sends RST to connected peers
- Host crash (with subsequent restart)
 - when restarted, TCP will respond to segments belonging to pre-crash connections with RST
- Host crash (without subsequent restart) or permanent network disconnection
 - connected peers trying to **send** will eventually timeout (can take several minutes)
 - programmers should use **their own** timeout mechanisms (or rely on the socket timeout mechanism offered by the Socket API)

Exercise

- Introduce mechanisms to avoid deadlocks in the Echo protocol implementations

Dependability through Redundancy: Process Groups


- **Process Group:** a group of processes perceived as a single, dependable process




- Implementing a Process Group requires distributed algorithms (protocols)
 - for managing the group (e.g., membership, leadership, data replication and consistency)
 - for making the group fault-tolerant

Process Group Organizations

Hierarchical
(primary-based protocols)



Flat
(replicated-write, active repl.
or quorum-based protocols)



Obtaining The Desired Fault-Tolerance

- A group is **k-fault tolerant** if it can withstand faults in k members
- The consensus protocols and the total number of processes necessary to achieve k-fault tolerance depend on the types of failures that are possible

Failure type	#Processes needed for k-fault tolerance	Consensus protocol Examples
crash, omission, timing, network	$k+1$	flooding
+incorrect response	$2k+1$	voting, Paxos
+arbitrary (byzantine)	$3k+1$	PBFT, Blockchain

The CAP (or Brewer's) Theorem

- Only 2 out of the following 3 properties can be obtained at the same time in a distributed system:
 - **C**onsistency
 - **A**vailability
 - **P**artition Tolerance
- Discussion
 - C, A, P are not on-off properties. The 2 out of 3 formulation may be misleading. A spectrum of alternatives is available.
- Readings:
 - E. Brewer. [CAP Twelve Years Later: How the “Rules” Have Changed](#), IEEE Computer, Vol. 45, No. 2, Feb. 2012, pp. 23-29.
 - P. Bailis, and K. Kingsbury. [The Network is Reliable](#), Communications of the ACM, Vol. 57, No. 9, Sept. 2014, pp. 48-55.